

# The Relational Data Borg is Learning

[fdbresearch.github.io](https://fdbresearch.github.io)

[relational.ai](https://relational.ai)

---

Dan Olteanu

University of Zurich

VLDB 2020 Keynote  
Virtual Tokyo, Sept 1, 2020



## Acknowledgments

FDB team, in particular:



Ahmet



Amir



Haozhe

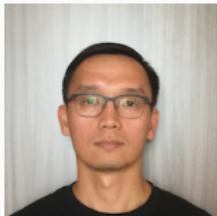


Max



Milos

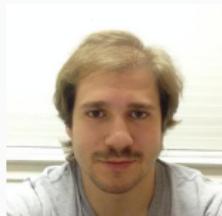
RelationalAI team, in particular:



Hung



Long



Mahmoud



Molham

# Database Research In Data Science Era

Reasons for DB research community to feel empowered:

1. **Pervasiveness of relational data** in data science
  - Hard fact
2. **Widespread need for efficient data processing**
  - Core to our community's *raison d'être*
3. **New processing challenges** posed by data science workloads
  - DB's approach reminiscent of Star Trek's Borg Collective

These reasons also serve as motivation for our work.

## Star Trek Borg

Co-opt technology and knowledge of alien species  
to become ever more powerful and versatile



# Relational Data Borg

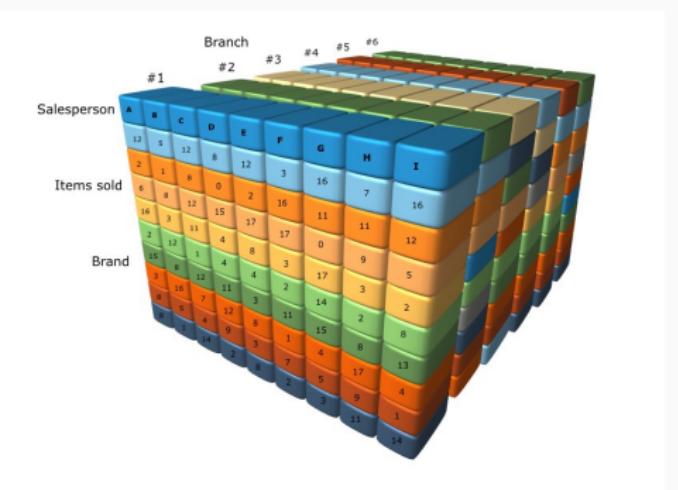
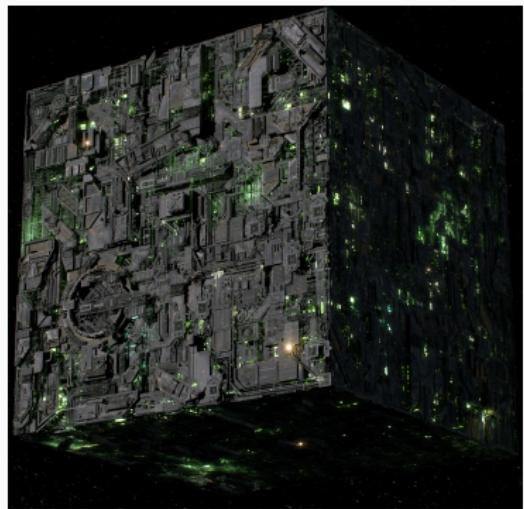


Assimilate ideas and applications of related fields  
to adapt to new requirements and  
become ever more powerful and versatile

Unlike in Star Trek, the Relational Data Borg

- moves fast
- has great skin complexion and
- is reasonably happy

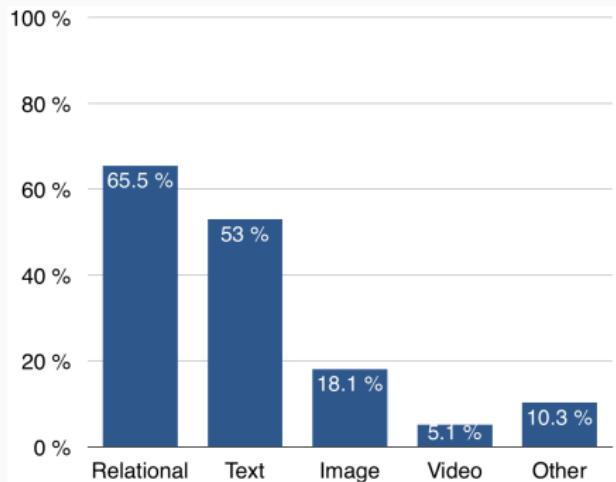
# Borg Cube vs Data Cube



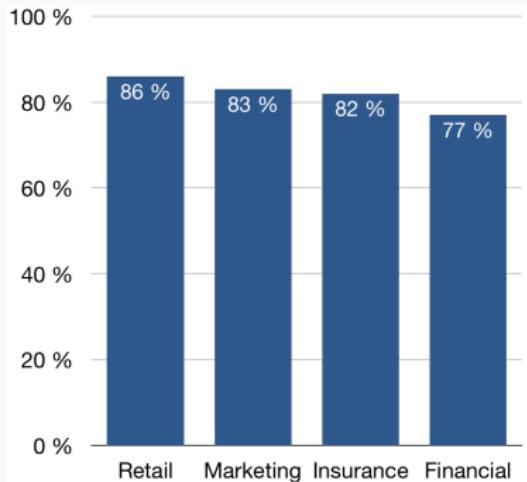


# Relational Data is Ubiquitous

**Kaggle Survey:** Most Data Scientists use Relational Data at Work!



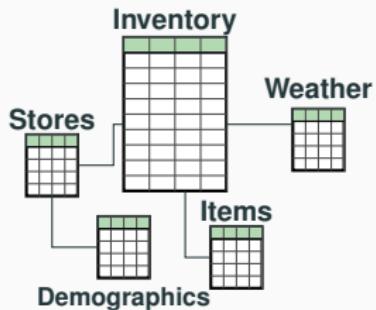
Overall



By Industry

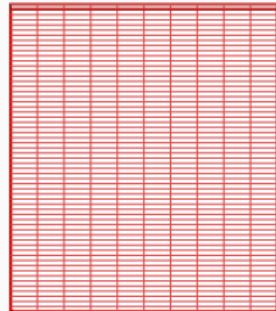
Source: The State of Data Science & Machine Learning 2017, Kaggle, October 2017  
(based on 2017 Kaggle survey of 16,000 ML practitioners)

# State of Affairs in Learning over Relational Data



Feature Extraction Query →  
Inventory  $\bowtie$  Stores  $\bowtie$  Items  
 $\bowtie$  Weather  $\bowtie$  Demographics

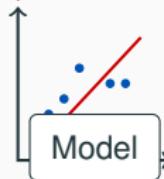
10,000s of Features



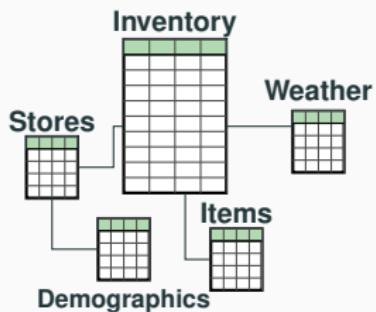
Training Dataset

ML Tool

Model



# State of Affairs in Learning over Relational Data

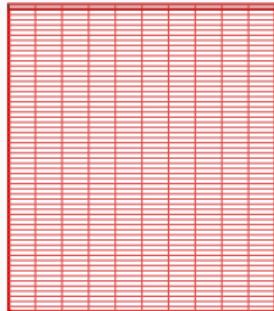


Feature Extraction Query

Inventory  $\bowtie$  Stores  $\bowtie$  Items  
 $\bowtie$  Weather  $\bowtie$  Demographics

Relational Data

10,000s of Features

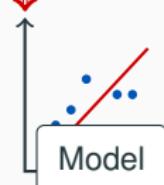


Training Dataset

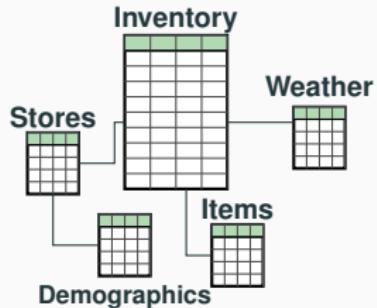
Structure-Agnostic Learning:

ML Tool

Model

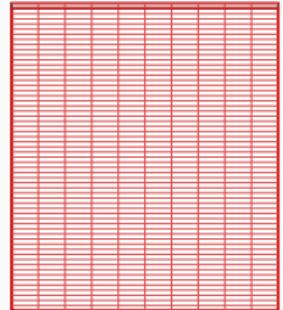


# State of Affairs in Learning over Relational Data



Feature Extraction Query →  
Inventory  $\bowtie$  Stores  $\bowtie$  Items  
 $\bowtie$  Weather  $\bowtie$  Demographics

10,000s of Features



Training Dataset

ML Tool

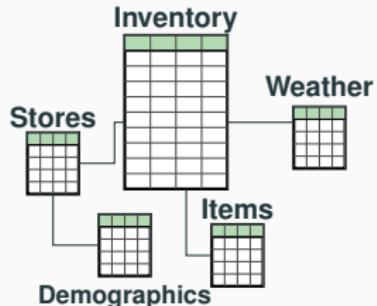
Model

Structure-Agnostic Learning:

1. Unnecessary data matrix materialization

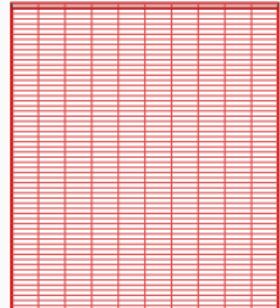
Relational structure carefully crafted by domain experts thrown away

# State of Affairs in Learning over Relational Data



Feature Extraction Query →  
Inventory  $\bowtie$  Stores  $\bowtie$  Items  
 $\bowtie$  Weather  $\bowtie$  Demographics

10,000s of Features



## Structure-Agnostic Learning:

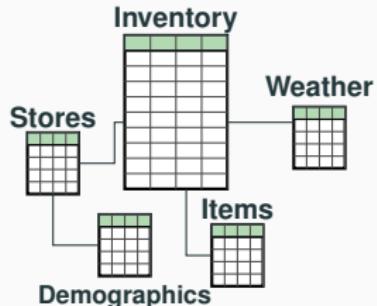
1. **Unnecessary** data matrix materialization

Relational structure carefully crafted by domain experts thrown away

2. **Expensive** data move

Training dataset can be **order-of-magnitude larger** than the input DB

# State of Affairs in Learning over Relational Data

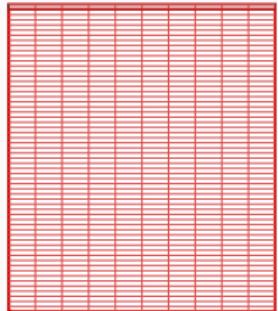


Feature Extraction Query

Inventory  $\bowtie$  Stores  $\bowtie$  Items  
 $\bowtie$  Weather  $\bowtie$  Demographics

Relational Data

10,000s of Features



Training Dataset

Structure-Agnostic Learning:

1. **Unnecessary** data matrix materialization

Relational structure carefully crafted by domain experts thrown away

2. **Expensive** data move

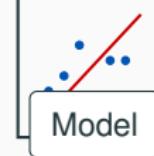
Training dataset can be **order-of-magnitude larger** than the input DB

3. **Bloated** one-hot encoding

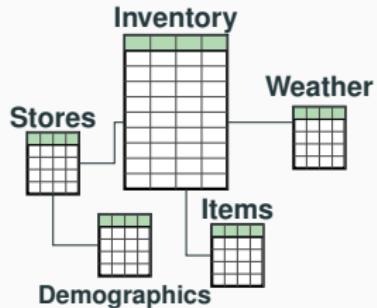
ML Tool



Model



# State of Affairs in Learning over Relational Data

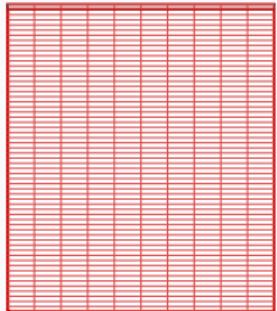


Feature Extraction Query

Inventory  $\bowtie$  Stores  $\bowtie$  Items  
 $\bowtie$  Weather  $\bowtie$  Demographics

Relational Data

10,000s of Features



Training Dataset

Structure-Agnostic Learning:

1. **Unnecessary** data matrix materialization

Relational structure carefully crafted by domain experts thrown away

2. **Expensive** data move

Training dataset can be **order-of-magnitude larger** than the input DB

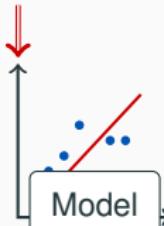
3. **Bloated** one-hot encoding

4. **High** maintenance cost

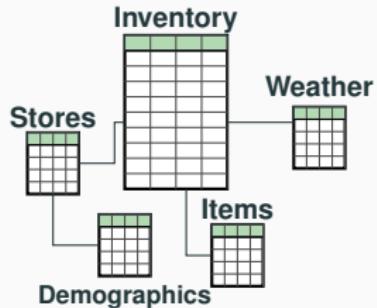
Recomputation from scratch after updates

ML Tool

Model



# State of Affairs in Learning over Relational Data

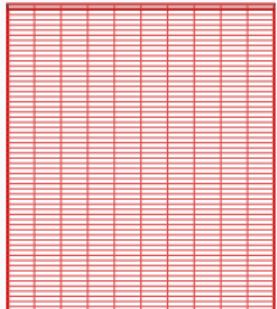


Feature Extraction Query

Inventory  $\bowtie$  Stores  $\bowtie$  Items  
 $\bowtie$  Weather  $\bowtie$  Demographics

Relational Data

10,000s of Features



Training Dataset

Structure-Agnostic Learning:

1. **Unnecessary** data matrix materialization

Relational structure carefully crafted by domain experts thrown away

2. **Expensive** data move

Training dataset can be **order-of-magnitude larger** than the input DB

3. **Bloated** one-hot encoding

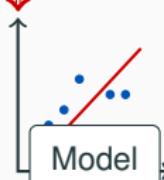
4. **High** maintenance cost

Recomputation from scratch after updates

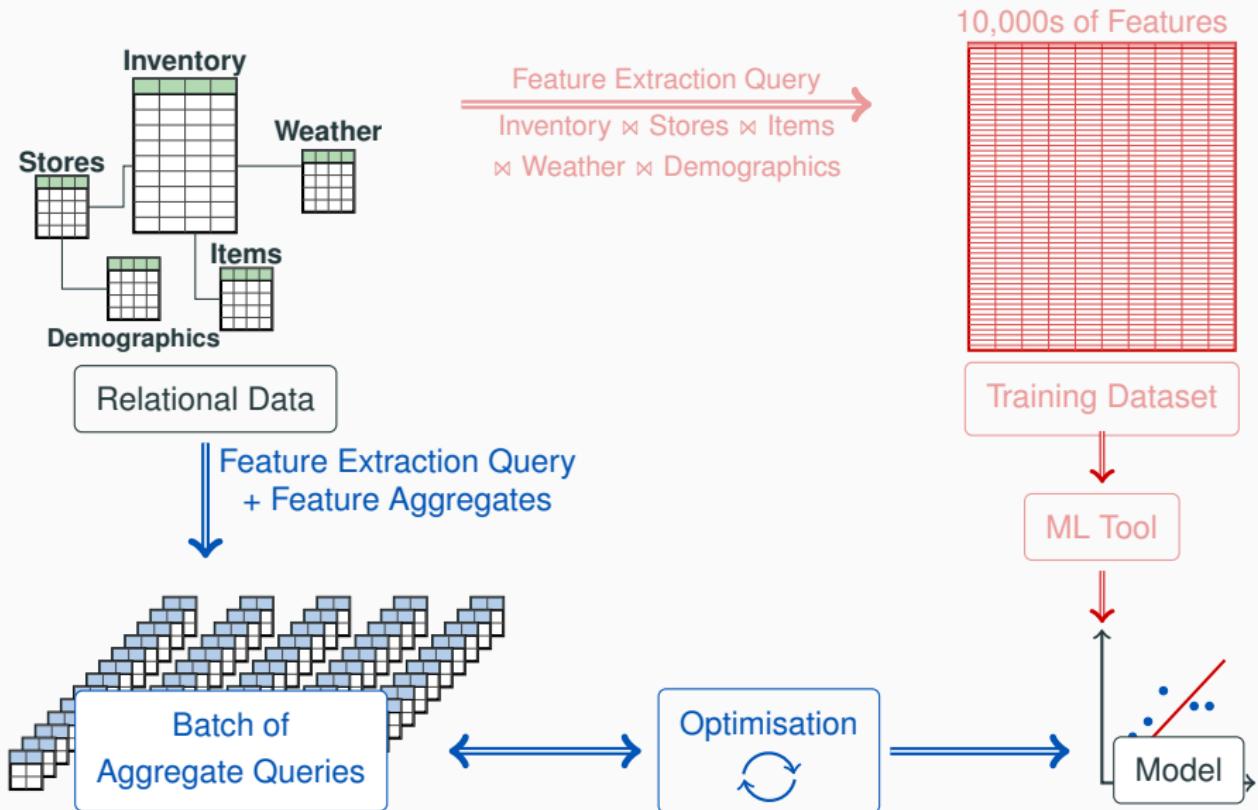
5. **Limitations** inherited from both DB and ML tools

ML Tool

Model



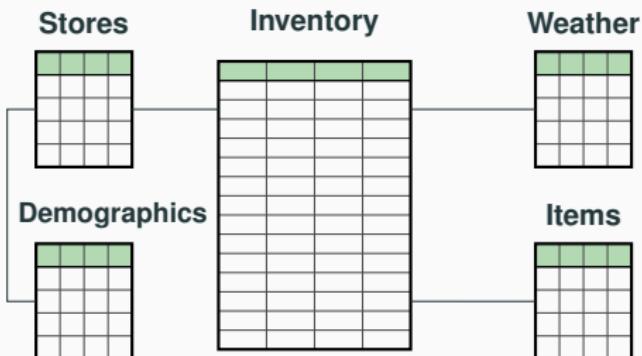
# Structure-Aware Learning over Relational Data



# Conjecture

The learning time and accuracy of the model can be drastically improved by exploiting the structure and semantics of the underlying multi-relational database.

## Structure-aware Learning FASTER than Feature Extraction Query alone



Relation	Cardinality	Arity (Keys+Values)	File Size (CSV)
Inventory	84,055,817	3 + 1	2 GB
Items	5,618	1 + 4	129 KB
Stores	1,317	1 + 14	139 KB
Demographics	1,302	1 + 15	161 KB
Weather	1,159,457	2 + 6	33 MB
Join	84,055,817	3 + 41	23GB

## Structure-aware versus Structure-agnostic Learning

Train a linear regression model to predict *inventory* given all features

### PostgreSQL+TensorFlow

	Time	Size (CSV)
Database	—	2.1 GB
Join	152.06 secs	23 GB
Export	351.76 secs	23 GB
Shuffling	5,488.73 secs	23 GB
Query batch	—	—
Grad Descent	7,249.58 secs	—
Total time	13,242.13 secs	

## Structure-aware versus Structure-agnostic Learning

Train a linear regression model to predict *inventory* given all features

	PostgreSQL+TensorFlow		Our approach (SIGMOD'19)	
	Time	Size (CSV)	Time	Size (CSV)
Database	–	2.1 GB	–	2.1 GB
Join	152.06 secs	23 GB	–	–
Export	351.76 secs	23 GB	–	–
Shuffling	5,488.73 secs	23 GB	–	–
Query batch	–	–	6.08 secs	37 KB
Grad Descent	7,249.58 secs	–	0.05 secs	–
Total time	13,242.13 secs		6.13 secs	

2,160× faster while being more accurate (RMSE on 2% test data)

## Structure-aware versus Structure-agnostic Learning

Train a linear regression model to predict *inventory* given all features

	PostgreSQL+TensorFlow		Our approach (SIGMOD'19)	
	Time	Size (CSV)	Time	Size (CSV)
Database	–	2.1 GB	–	2.1 GB
Join	152.06 secs	23 GB	–	–
Export	351.76 secs	23 GB	–	–
Shuffling	5,488.73 secs	23 GB	–	–
Query batch	–	–	6.08 secs	37 KB
Grad Descent	7,249.58 secs	–	0.05 secs	–
Total time	13,242.13 secs		6.13 secs	

2,160× faster while being more accurate (RMSE on 2% test data)

TensorFlow trains one model. Our approach takes < 0.1 sec for any extra model over a subset of the given feature set.

# TensorFlow's Behaviour is the Rule, not the Exception!

Similar behaviour (or outright failure) for more:

- **datasets**: Favorita, TPC-DS, Yelp, Housing
- **systems**:
  - used in industry: R, scikit-learn, Python StatsModels, mpack, XGBoost, MADlib
  - academic prototypes: Morpheus, libFM
- **models**: decision trees, factorisation machines,  $k$ -means, ..

This is to be contrasted with the scalability of DBMSs!

**How to achieve this performance  
improvement?**

## Idea 1: Turn the ML Problem into a DB Problem



## Through DB Glasses, Everything is a Batch of Queries

Workload	Query Batch
Linear Regression	$\text{SUM}(X_i * X_j)$
Covariance Matrix	$\text{SUM}(X_i) \text{ GROUP BY } X_j$ $\text{SUM}(1) \text{ GROUP BY } X_i, X_j$
Decision Tree Node	$\text{VARIANCE}(Y) \text{ WHERE } X_j = c_j$
Mutual Information	$\text{SUM}(1) \text{ GROUP BY } X_i$
R $k$ -means	$\text{SUM}(1) \text{ GROUP BY } X_j$ $\text{SUM}(1) \text{ GROUP BY Center}_1, \dots, \text{Center}_k$

## Through DB Glasses, Everything is a Batch of Queries

Workload	Query Batch
Linear Regression	$\text{SUM}(X_i * X_j)$ [ WHERE $\sum_k X_k * w_k < c$ ]
Covariance Matrix	$\text{SUM}(X_i)$ GROUP BY $X_j$ [ WHERE ... ]
(Non)poly. loss	$\text{SUM}(1)$ GROUP BY $X_i, X_j$ [ WHERE ... ]
Decision Tree Node	$\text{VARIANCE}(Y)$ WHERE $X_j = c_j$
Mutual Information	$\text{SUM}(1)$ GROUP BY $X_i$
R $k$ -means	$\text{SUM}(1)$ GROUP BY $X_j$ $\text{SUM}(1)$ GROUP BY Center <sub>1</sub> , ..., Center <sub><math>k</math></sub>

## Through DB Glasses, Everything is a Batch of Queries

Workload	Query Batch	# Queries
Linear Regression	$\text{SUM}(X_i * X_j) \quad [\text{WHERE } \sum_k X_k * w_k < c]$	814
Covariance Matrix	$\text{SUM}(X_i) \text{ GROUP BY } X_j \quad [\text{WHERE ...}]$	
(Non)poly. loss	$\text{SUM}(1) \text{ GROUP BY } X_i, X_j \quad [\text{WHERE ...}]$	
Decision Tree Node	$\text{VARIANCE}(Y) \text{ WHERE } X_j = c_j$	3,141
Mutual Information	$\text{SUM}(1) \text{ GROUP BY } X_i$	56
R <sup>k</sup> -means	$\text{SUM}(1) \text{ GROUP BY } X_j$ $\text{SUM}(1) \text{ GROUP BY Center}_1, \dots, \text{Center}_k$	41

(# Queries shown for Retailer dataset with 39 attributes)

Queries in a batch:

- Same aggregates but over different attributes
- Expressed over the same join of the database relations

**AMPLE** opportunities for sharing computation in a batch.

# Models under Consideration

So far:

- Polynomial regression
- Factorisation machines
- Classification/regression trees
- Mutual information
- Chow Liu trees
- $k$ -means clustering
- $k$ -nearest neighbours
- (robust, ordinal) PCA
- SVM

On-going:

- Boosting regression trees
- AdaBoost
- Sum-product networks
- Random forests
- Logistic regression
- Linear algebra:
  - QR decomposition
  - SVD
  - low-rank matrix factorisation

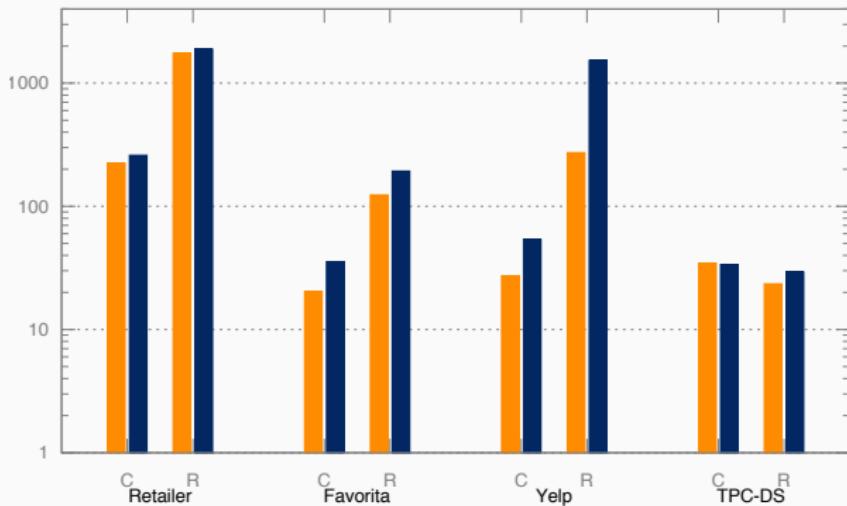
All these cases can benefit from **structure-aware computation**

## **Natural Attempt:**

Use Existing DB System to Compute Query Batch

# Existing DBMSs are **NOT** Designed for Query Batches

Relative Speedup for **Our Approach** over **DBX** and **MonetDB**

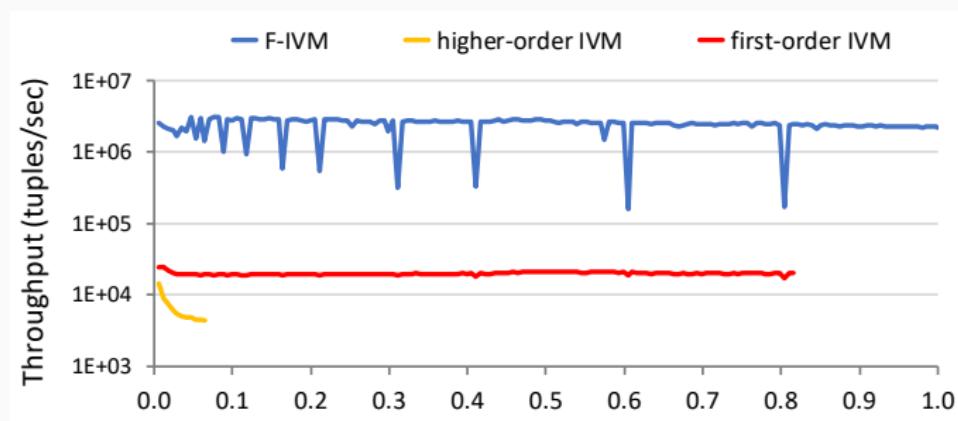


C = Covariance Matrix; R = Regression Tree Node; AWS d2.xlarge (4 vCPUs, 32GB)

# Existing DSMSs are **NOT** Designed for Query Batches

Task: Maintain the covariance matrix over Retailer

- Round-robin insertions in all relations
- All maintenance strategies implemented in DBToaster



Azure DS14, Intel Xeon, 2.40GHz, 112GB, 1 thread; one hour timeout